Safety-Aware Vehicle Routing for Quick-Commerce Delivery in Indian Cities

An ε-Constraint Formulation with Soft Time Windows

Pranav Bhadane
Roll No. 21IM30017

Department of Industrial and Systems Engineering
Indian Institute of Technology Kharagpur

Supervisor: Prof. Subhajit Sidhanta

M.Tech. Thesis Presentation, 2026

Introduction: The Q-Commerce Routing Problem

Setting

- Platforms like **Blinkit**, Zepto, Instamart promise **10-minute delivery** from neighbourhood dark stores
- Fleet of **two-wheeler riders** in dense Indian cities
- High dispatch frequency, narrow residential streets

The Safety Problem

- India: **4,61,312 road accidents** and **1,68,491 fatalities** in 2022 (MoRTH)
- Fastest route $\neq$ safest route
- Conventional VRP treats safety as an

Why Existing Models Fail

- Minimise only travel time or distance
- Crash risk as a weighted term $\Rightarrow$ **hidden trade-offs**
- No explicit residential-street protection
- No calibration story for risk estimates

Literature Review

VRPTW with Soft Time Windows

- Solomon (1987) — benchmark VRPTW instances
- Koskosidis et al. (1992) — soft TW with linear penalty
- Taillard et al. (1997) — convex penalties rarely linearised
- Desrochers et al. (1992) — exact column generation

Safety-Aware Routing

- Hazmat origins: Zografos & Erkut
- Hoseinzadeh et al. (2020) — cumulative risk, *weighted-sum* objective

Quick-Commerce & Indian Context

- Sinchaisri et al. (2023) — SLA pressure in q-commerce
- Das et al. (2020) — Indian gig-rider compliance

Metaheuristics

- Prins (2004) — giant-tour + Bellman–Ford split
- Ropke & Pisinger (2006) — Adaptive LNS

Data & Calibration

- Boeing/OSMnx (2017) — street networks

Research Gaps

Gap 1
No framework combining safety constraints + soft TW + q-commerce SLA

Gap 2
Crash risk uncalibrated — no reproducible public-source chain for India

Gap 3
Solver comparisons use different evaluation code (apples-to-pears bias)

Gap 4
Residential-street penetration not modelled as an explicit operator knob

Gap 5
No benchmark on real Indian dark-store coordinates — synthetic depots used

$\Downarrow$

This thesis addresses all five gaps in one integrated, reproducible framework.

Contributions

1. **SA-VRPTW Formulation** — crash exposure as hard ε-constraint ($\sum R_{ij} x_{ij}^k \leq \bar{R}$); per-route + fleet-wide residential caps $\bar{H}_k = 100$
2. **BASM Calibration** — reproducible public-source chain (MoRTH 2022 + Mohan et al. 2017 + OSM proxy) without iRAD microdata; explicit π_{ij} formula published
3. **Unified Solver Contract** — MILP, GA, ALNS all share identical F_1 evaluation and feasibility checker
4. **Key Empirical Finding** — GA with BF-split outperforms ALNS at $N \geq 100$ in quality AND speed (**325**$\times$ **faster at** $N = 200$)
5. **Five-City Benchmark** — real Blinkit dark-store coordinates (Bengaluru, Delhi, Gurugram, Mumbai, Pune); 428 feasible rows
6. **Fully Reproducible Pipeline** — Hydra config, seeded runs, manifest-checkpointed grid, MIT licence

Problem Formulation: SA-VRPTW

Primary Objective (minimise service-quality cost):

$$F_1 = \sum_k \sum_{i \in C} \Big[w_e \, w_i^k + \big(e^{\beta_{\text{stw}} \tau_i^k} - 1 \big) \Big]$$

- w_i^k: early-arrival idle wait
- τ_i^k: lateness beyond 10-min ETA
- Convex penalty $\Rightarrow$ realistic SLA cost

ε-**Constraints** (hard safety bounds):

$$\textstyle\sum_{(i,j)} R_{ij} \, x_{ij}^k \le \bar{R} \quad \forall k \;\; \text{(crash survival)}$$

$$\textstyle\sum_{(i,j)} H_{ij} \, x_{ij}^k \le \bar{H}_k \quad \forall k \;\; \text{(residential cap)}$$

Super-Graph Arc Quantities

Shortest-time path P_{uv} on street graph:

$$T_{uv} = \sum_{e \in P} t_e, \quad R_{uv} = \sum_{e} -\ln(1 - r_e), \quad H_{uv} = \sum_{e}$$

Key Constants

$Q = 2$ items	$T_{\max} = 35$ min
$\bar{H}_k = 100$ edges	$\beta_{\text{stw}} = 0.12$
PWL breakpoints:	$\{0, 2, 5, 10, 15, 20, 30\}$

Multi-depot coupling, MTZ sub-tour elimination, and arrival-time linking constraints fully stated in Chapter 3.

Data Pipeline and BASM Calibration

OSMnx street graph (drive, bbox)

↓

BPR Congestion: $t_{ij} = t_{ij}^{free}(1 + 0.15(V/C)^4)$

↓

BASM Risk: $r_{ij} = \mathrm{clip}(\lambda \cdot \ell \cdot \sigma \cdot \pi_{ij}, 0, 0.99)$

↓

Blinkit KML → k-means depots

↓

Residential-weighted

BASM Local Proxy:

$$\pi_{ij} = \mathrm{clip}\big(1 + 0.40(2\beta_{ij} - 1) + 0.30 s_{ij} + 0.30 c_{ij},\ 0.5$$

- β_{ij}: normalised edge betweenness
- s_{ij}: local signal density
- c_{ij}: crossing density

Calibration chain:

- MoRTH 2022 — category-level fatality totals
- Mohan et al. 2017 — arterial > collector > local
- OSM proxy — **no iRAD microdata**

Solver Architecture: Unified Evaluation Contract

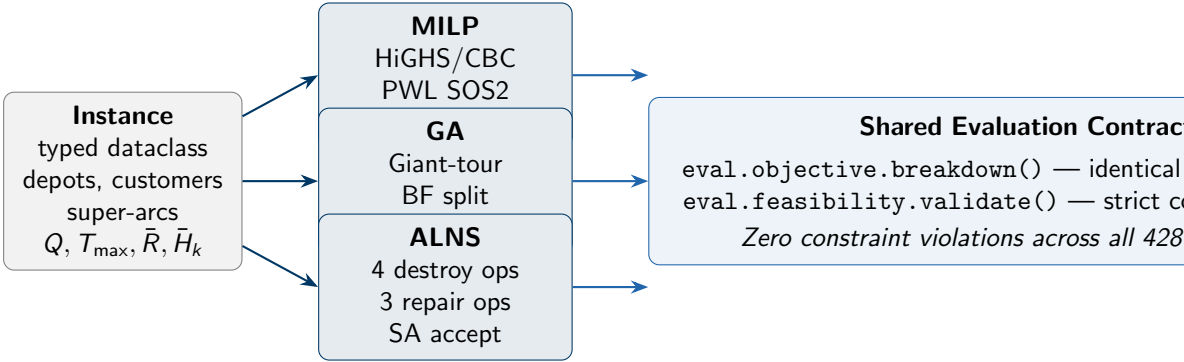

MILP: exact benchmark; PWL outer-envelope for $e^{\beta\tau} - 1$ via SOS2

GA: OX1 crossover; swap/insert/reverse mutation; BF split enforces hard caps

ALNS: Shaw/worst/random/risk-cluster destroy; greedy & regret-2/3

Algorithm Design: GA + ALNS Pseudocode

Algorithm 1 — GA with Bellman–Ford Split

```
Init P random customer permutations
Decode via BF split; evaluate F_1 + penalty
for g = 1 to G = 200 do
    Carry top-e = 4 elites forward
    p_1, p_2 ← tournament(k = 5)
    Child ← OX1 crossover (p_x = 0.85)
    Apply swap / insert / reverse (p_m = 0.20)
    Decode; evaluate; sort; truncate to P
Validate incumbent via validate()
return s*
```

BF split adds 10^9 sentinel for routes violating H_k, $\bar{R}$, or $T_{\max}$ — hard constraints never silently broken.

Algorithm 2 — ALNS

```
s_0 ← GreedyConstruct(I)
T_SA ← 0.05 · F_1(s_0)
for t = 1 to 5000 do
    q ← ⌈n · U(0.15, 0.45)⌉
    D ← destroy op ∝ w^d
    R ← repair op ∝ w^r
    s' ← R(D(s, q))
    Accept/reject via SA (η = 0.99975)
    Update σ = (33, 9, 1, 0) weights
Validate; return s*
```

Destroy: random, worst-cost, Shaw relatedness, risk-cluster.

Repair: greedy insert, regret-2, regret-3.

Code Snapshot: Bellman–Ford Split with Hard-Constraint Sentinels

```python
# src/savrptw/solvers/_common.py
HARD_VIOLATION = 1e9

for i in range(n):
    cumulative_load = 0
    for j in range(i + 1, n + 1):
        cumulative_load += customers[j-1].demand
        if cumulative_load > instance.Q:
            break  # capacity prune (O(n^2))

        rb = build_route(instance, depot,
                         customers[i:j],
                         penalty_weights=pw)
        # Soft F1 + penalty mass
        route_cost = rb.f1_contribution \
                   + rb.violation_mass

        # Hard-constraint sentinels
        if rb.total_H > instance.H_cap_route:
            route_cost += HARD_VIOLATION
```

Why this design?

- Shared by GA, ALNS, and MILP evaluator
- F_1 identical across all three — no apples-to-pears bias
- 10^9 sentinel: BF only selects a violating arc when *no feasible split exists*
- $O(n^2)$ — fast: 26 s at $N = 200$

Project layout:

```
src/savrptw/
  solvers/ — milp, ga, alns
  instance/ — generator, super_arc
  risk/ — basm calibration
```

Experimental Setup

Cities and Dark-Store Coverage

City	Blinkit stores
Bengaluru	113
Delhi	111
Mumbai	96
Pune	39
Gurugram	38
Hyderabad	0 (excluded)

Grid:

- $N \in \{20, 50, 100, 200\}$ customers
- Depots: $\{3, 5, 5, 8\}$ per N
- Up to 10 random seeds per cell
- Solvers: GA and ALNS

Reproducibility

- Hydra config composition
- Manifest-checkpointed grid runner — crash-safe resumption
- All parameters in `conf/*.yaml`
- BASM guard: refuses uncalibrated source

Evaluation

- Strict `validate()` after every solve
- Crash Monte Carlo: 10,000 Bernoulli trials
- Beta-prior compliance sweep:

Results: F_1 vs Instance Size (All Cities)

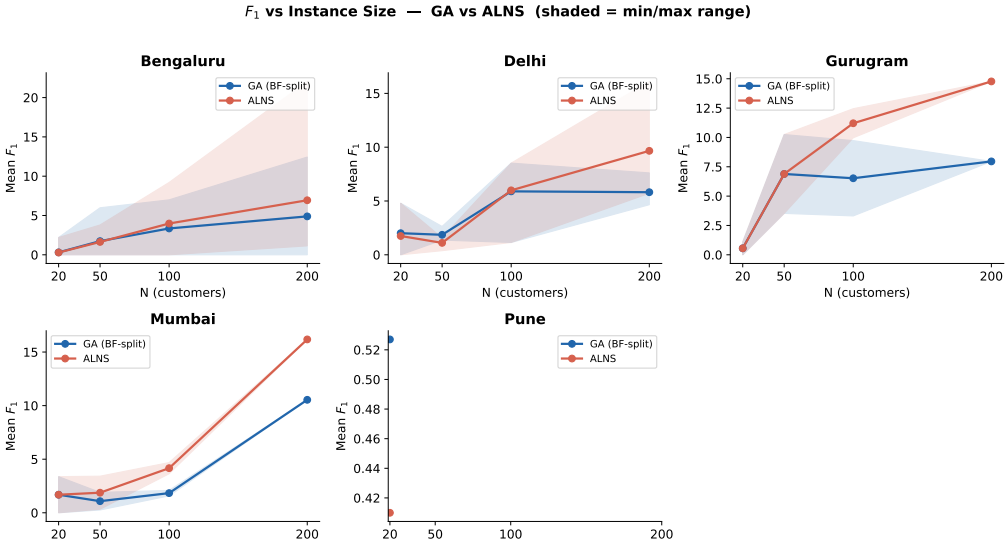

Results: Wall-Clock Time and Speed-Up Ratio

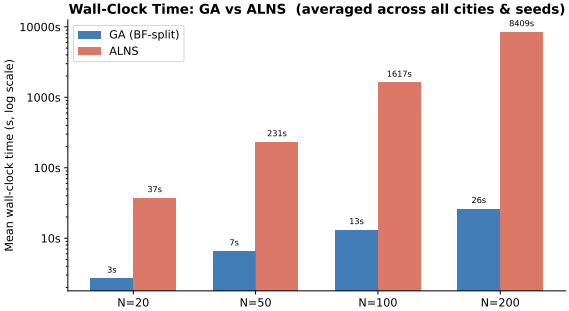

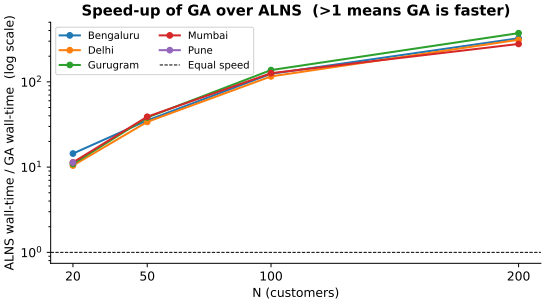

Solver	$N = 20$	$N = 50$	$N = 100$	$N = 200$
GA	2.7 s	6.5 s	13 s	**26 s**
ALNS	31 s	213 s	1669 s	**8407 s**

GA is $\sim$325$\times$ faster at $N = 200$ and produces lower F_1

Results: City-Level Analysis

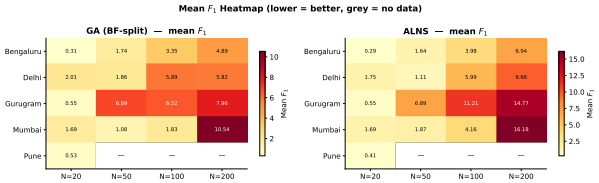

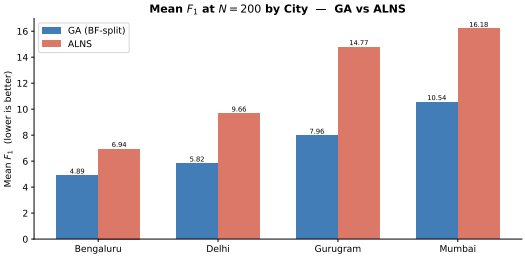

- **Bengaluru** (113 stores) — most stable; lowest F_1; recommended for sensitivity studies
- **Gurugram & Mumbai** — elevated F_1 at $N \geq 50$ due to fewer depot clusters
- **GA $\leq$ ALNS** in F_1 across all cities at $N = 200$ — consistent
- **Pune** — only $N = 20$ in current run; insufficient to generalise

Limitations & Future Work

Honest Limitations

- BASM is a *calibrated prior*, not a measurement — no iRAD microdata
- Super-graph fixes paths: optimisation over customer **orderings**, not physical street choices
- MILP scales to $N \leq 50$ only (CBC); commercial solver needed for larger
- Pareto $\bar{R}$–$\bar{H}$ sweep deferred to next run
- Behavioural compliance from single field-data study

Future Directions

1. **Risk-aware path selection** — jointly optimise orderings + physical streets
2. **Pareto sweep** — $\bar{R} \times \bar{H}$ grid traces safety/efficiency frontier
3. **Real-time congestion** — replace static BPR with live feeds
4. **Stochastic demand** — non-stationary order arrivals
5. **Behavioural meta-analysis** — multi-study compliance prior
6. **Direct crash microdata** — upgrade BASM if iRAD access obtained

Conclusion

What We Delivered

- SA-VRPTW formulation with ε-constraints for crash survival and residential-road use
- Reproducible public-source BASM calibration for India — no restricted data needed
- MILP, GA, and ALNS under a unified F_1/feasibility contract
- Empirical benchmark: 428 feasible rows, 5 cities, 4 problem sizes

Key Takeaways

- **ε-constraint framing** gives operators an explicit safety dial — no hidden weight tuning
- **GA with Bellman–Ford split** is recommended: *both* faster and higher-quality than ALNS at $N \geq 100$
- **Calibration honesty** is achievable without restricted microdata — chain is fully reproducible
- Full pipeline released under **MIT licence**: `github.com/Coderopp/MTP-II`

Thank You

Questions & Discussion

Author: Pranav Bhadane (21IM30017)
Supervisor: Prof. Subhajit Sidhanta
Department: Industrial & Systems Engg., IIT Kharagpur
Code: `github.com/Coderopp/MTP-II`

Key references: Solomon (1987) · Prins (2004) · Ropke & Pisinger (2006)
MoRTH (2022) · Mohan et al. (2017) · Boeing/OSMnx (2017)